

```
In [ ]: #GitHub link for this assignment - https://github.com/ass2testing/week4.git
```

Step 1: Import Necessary Libraries First, import the libraries you'll need for data manipulation, visualization, and model building.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import nltk
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Step 2: Load the Data Load the training and test data from the Kaggle dataset

```
In [2]: train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

Step 3: Explore the Data Get a sense of what the data looks like.

```
In [3]: print(train_data.head())
print(train_data.info())
print(train_data.describe())
```

```

      id keyword location text \
0    1    NaN      NaN Our Deeds are the Reason of this #earthquake M...
1    4    NaN      NaN Forest fire near La Ronge Sask. Canada
2    5    NaN      NaN All residents asked to 'shelter in place' are ...
3    6    NaN      NaN 13,000 people receive #wildfires evacuation or...
4    7    NaN      NaN Just got sent this photo from Ruby #Alaska as ...

```

```

target
0    1
1    1
2    1
3    1
4    1

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7613 entries, 0 to 7612
```

```
Data columns (total 5 columns):
```

```

#   Column      Non-Null Count  Dtype
---  -----  -
0   id        7613 non-null    int64
1   keyword    7552 non-null    object
2   location   5080 non-null    object
3   text       7613 non-null    object
4   target     7613 non-null    int64

```

```
dtypes: int64(2), object(3)
```

```
memory usage: 297.5+ KB
```

```
None
```

```

              id      target
count  7613.000000  7613.000000
mean    5441.934848    0.42966
std     3137.116090    0.49506
min       1.000000    0.00000
25%    2734.000000    0.00000
50%    5408.000000    0.00000
75%    8146.000000    1.00000
max   10873.000000    1.00000

```

Step 4: Preprocess the Text Data Clean and preprocess the text data by removing stopwords, punctuation, and applying other necessary transformations.

```

In [4]: nltk.download('stopwords')
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    text = text.lower() # Convert to Lowercase
    text = re.sub(r'\d+', '', text) # Remove numbers
    text = re.sub(r'https?://\S+|www\.\S+', '', text) # Remove URLs
    text = re.sub(r'^a-zA-Z\s', '', text) # Remove punctuation
    text = text.split()
    text = [word for word in text if word not in stop_words]
    text = ' '.join(text)
    return text

train_data['cleaned_text'] = train_data['text'].apply(preprocess_text)
test_data['cleaned_text'] = test_data['text'].apply(preprocess_text)

```

```

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Gyadav\AppData\Roaming\nltk_data...
[nltk_data] Unzipping corpora\stopwords.zip.

```

Step 5: Vectorize the Text Data Convert the text data into numerical features using TF-IDF Vectorizer.

```
In [5]: vectorizer = TfidfVectorizer(max_features=1000)
X = vectorizer.fit_transform(train_data['cleaned_text']).toarray()
y = train_data['target']

X_test = vectorizer.transform(test_data['cleaned_text']).toarray()
```

Step 6: Train-Test Split Split the training data into training and validation sets.

```
In [6]: X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 7: Build and Train the Model Use a simple Logistic Regression model to start.

```
In [7]: model = LogisticRegression()
model.fit(X_train, y_train)
```

```
Out[7]: ▾ LogisticRegression
LogisticRegression()
```

Step 8: Evaluate the Model Check the accuracy and other metrics on the validation set.

```
In [9]: y_pred = model.predict(X_val)
print('Accuracy:', accuracy_score(y_val, y_pred))
print('Confusion Matrix:\n', confusion_matrix(y_val, y_pred))
print('Classification Report:\n', classification_report(y_val, y_pred))
```

Accuracy: 0.7984241628365069

Confusion Matrix:

```
[[766 108]
 [199 450]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.88	0.83	874
1	0.81	0.69	0.75	649
accuracy			0.80	1523
macro avg	0.80	0.78	0.79	1523
weighted avg	0.80	0.80	0.80	1523

Step 9: Make Predictions on the Test Set Finally, use the trained model to make predictions on the test set.

```
In [10]: test_predictions = model.predict(X_test)
submission = pd.DataFrame({'id': test_data['id'], 'target': test_predictions})
submission.to_csv('submission.csv', index=False)
```

```
In [ ]:
```